

RAM CIRIGIRI

Austin, TX | +1 321-576-7292 | ram.cirigiri@gmail.com | linkedin.com/in/rciri | rcirigiri.github.io

Engineering Leader, AI Systems

PROFESSIONAL SUMMARY

Engineering leader with 15+ years building high-scale distributed systems and production-grade AI platforms. Proven track record managing cross-functional teams (8+ engineers) at Fortune 25 companies delivering \$5M+ in annual business impact. Specialist in real-time voice AI, multi-agent orchestration, and RAG pipelines — from research to production with strict quality, compliance, and cost-efficiency standards.

TECHNICAL CORE

Engineering Leadership	0 → 7 Team Building · Promotions · Technical Roadmap Ownership · Cross-functional Delivery · Engineering Standards
AI Architecture & Agents	LangGraph · Multi-agent Systems · RAG · Hybrid BM25+Vector · Prompt Engineering · Vector Search · Cosmos DB
AI Quality & Compliance	Eval Suites (CI-gated) · LLM Audit Logging · Hallucination Guard · PII Redaction · Feedback-to-Golden-Dataset
Voice & Real-time	Deepgram · ElevenLabs · Sarvam AI · Dual-STT Arbitration · WebSocket Streaming · Barge-In · High-Concurrency
Cloud & Infrastructure	Azure · Cosmos DB · Event Hubs · Python · FastAPI · C#.NET · Redis · Kafka · Kubernetes · Docker · Terraform

Technologies: LangGraph · Claude API · OpenAI API · Azure · Cosmos DB · Redis · Kafka · FastAPI · Deepgram · Sarvam AI · Docker · Kubernetes · Terraform · Next.js · WebSocket · RAG · BM25 · Vector Search

REPRESENTATIVE EXPERIENCE

Cognitive Commerce Labs (CCL)

Founder

June 2025 – Present

Austin, TX

Founder of an AI commerce startup; built two conversational-AI products end-to-end — a three-channel restaurant ordering system (voice, chat, and phone), live in production across 5 restaurants, and a grocery product-discovery system in pilot — owning architecture, delivery, reliability, and a small cross-functional team.

- Built a real-time, three-channel conversational ordering system (voice, chat, and phone); a patent-pending dual-engine speech-arbitration layer (USPTO #64/026,827) selects the more accurate transcript in ~1ms, reducing mis-heard orders on accented and multilingual speech.
- Architected the ordering agent as an 18-node LangGraph state machine with 32 grounded tools, completing most orders in a single exchange; a deterministic fast-path and tiered model routing (~70% of turns on a low-cost model) cut ~1s of latency per turn and lowered serving cost.
- Built an AI product-discovery system architected for grocery catalogs of 2M+ SKUs with a tiered resolution pipeline (deterministic matching → embeddings → LLM for only the ~3% long tail) and an ML entity-resolution flywheel reaching 95%+ automated resolution — avoiding an estimated ~40 hours/week of manual reconciliation.
- Raised search quality to 85% Precision@10 (+15–20%) with hybrid BM25 + vector retrieval (RRF fusion), and cut P95 latency 800ms → 150ms via async execution and query optimization — load-tested to 100+ QPS at 99.9% availability.
- Kept infrastructure cost to ~\$500/mo — about 1% of a naive LLM-per-query design — via a two-tier cache (80%+ hit rate); delivered the system greenfield to production in 6 months.
- Engineered for production reliability and compliance: speech- and LLM-provider failover, per-session distributed locking, PII redaction and audit logging on every model call, p95/error-rate observability, and a 630-test CI-gated regression suite.
- Led a cross-functional team of 2–4 engineers across product, data science, and infrastructure; designed the canonical-normalization architecture to scale to 10,000 stores.

USPTO Provisional Patent (2026):

→ Domain-Aware Multilingual Speech Recognition Arbitration · #64/026,827

General Motors (Fortune 25)

Software Engineering Manager

2022 – May 2025

Austin, TX

- People Leadership** Grew team 0 → 7 (5 FTE + 2 offshore) for a greenfield claims platform targeting a \$300B insurance market. Promoted Senior → Tech Lead, developed Mid → Senior through module ownership; navigated one PIP to resolution. Architectural patterns adopted as the global reference implementation. Engineers developed under this model have subsequently moved into senior technical leadership roles.
- AI & Compliance** Led Claims Document Intelligence Tool — HIPAA-compliant architecture with PII/PHI redaction and customer-managed encryption keys, establishing the first safe GenAI use for sensitive insurance records at GM. Established the architectural and commercial foundation for \$2M+ in projected annual savings — validated with business and compliance stakeholders.

Publix Super Markets (Fortune 100)

2019 – 2022

Technical Programme Lead & Lead Architect

Lakeland, FL · 6-engineer cross-functional team · 2 concurrent programmes

- **Sales360** Architected a high-throughput event-driven platform (on-prem Kafka → Azure Event Hubs) as single source of truth for 1,200+ stores. Eliminated a 23-hour data staleness window via real-time streaming. Zero incidents at production cutover.
- Designed a two-tier caching architecture (in-memory L1 + Redis L2) for Avro schema resolution in an Azure Function processing 5M+ daily Kafka transactions at 500 TPS.
- **Central Refunds** Real-time compliance and fraud prevention platform · sub-6s SLA · configurable rules engine · full audit trail · \$3M annual savings.

American Express (Fortune 100)

2013 – 2019

Technical Lead — IVR & Customer Service Platform

Clearwater, FL · Team of 5 · Technology Champion

- Led Amex Serve platform processing \$5.5M daily transactions (peak \$20M). Designed intelligent IVR routing that reduced live agent transfers, saving \$500K+ annually.

Inmarsat (Satellite Communications)

2009 – 2013

Senior Software Engineer

Melbourne, FL

- Developed mission-critical satellite communications software for global maritime and aviation customers. Deep foundation in high-reliability distributed systems and fault-tolerant real-time data processing.

EDUCATION

MBA

University of Florida — Warrington College of Business

MS in Computer Science

Florida Institute of Technology — Distributed Systems & Software Engineering